

## Spring 2023: CS5710 – Machine Learning

### In-Class Programming Assignment-4

GitHub Link - [https://github.com/raimukul/MachineLearning\\_Assignments](https://github.com/raimukul/MachineLearning_Assignments)

Video link-

<https://drive.google.com/file/d/12RFhEuUN1nxvux8EN39QYeSSReYUybeZ/view?usp=sharing>

Code:

#### 1. Pandas

1. Read the provided CSV file 'data.csv'.

<https://drive.google.com/drive/folders/1h8C3mLsso-R-sIOLsvoYwPLzy2fJ4IOF?usp=sharing>

2. Show the basic statistical description of the data.
3. Check if the data has null values. a. Replace the null values with the mean.
4. Select at least two columns and aggregate the data using: min, max, count, mean.
5. Filter the data frame to select the rows with calories values between 500 and 1000.
6. Filter the data frame to select the rows with calories values > 500 and pulse < 100.
7. Create a new "df\_modified" data frame that contains all the columns from df except for "Maxpulse".
8. Delete the "Maxpulse" column from the main df dataframe
9. Convert the datatype of Calories column to int datatype.
10. Using pandas create a scatter plot for the two columns (Duration and Calories).

In [1]:

```
#Read the provided CSV file 'data.csv'.  
https://drive.google.com/drive/folders/1h8C3mLsso-R-sIOLsvoYwPLzy2fJ4IOF?usp=sharing
```

```
import pandas as pd  
  
df = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/data/data.csv')
```

In [17]:

```
print(df)
```

	Duration	Pulse	Maxpulse	Calories
0	60	110	130	409.1
1	60	117	145	479.0
2	60	103	135	340.0
3	45	109	175	282.4
4	45	117	148	406.0
..	...	...	...	...
164	60	105	140	290.8
165	60	110	145	300.0
166	60	115	145	310.2
167	75	120	150	320.4
168	75	125	150	330.4

[169 rows x 4 columns]

In [2]:

```
df = pd.DataFrame(df)
```

In [3]:

*#Show the basic statistical description about the data.*

```
df=df.describe()
```

```
df
```

Out[3]:

	Duration	Pulse	Maxpulse	Calories
<b>count</b>	169.000000	169.000000	169.000000	164.000000
<b>mean</b>	63.846154	107.461538	134.047337	375.790244
<b>std</b>	42.299949	14.510259	16.450434	266.379919
<b>min</b>	15.000000	80.000000	100.000000	50.300000
<b>25%</b>	45.000000	100.000000	124.000000	250.925000

	Duration	Pulse	Maxpulse	Calories
<b>50%</b>	60.000000	105.000000	131.000000	318.600000
<b>75%</b>	60.000000	111.000000	141.000000	387.600000
<b>max</b>	300.000000	159.000000	184.000000	1860.400000

In [4]:

```
#Check if the data has null values.
df = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/data/data.csv')
df.isnull()
```

Out[4]:

	Duration	Pulse	Maxpulse	Calories
<b>0</b>	False	False	False	False
<b>1</b>	False	False	False	False
<b>2</b>	False	False	False	False
<b>3</b>	False	False	False	False
<b>4</b>	False	False	False	False
<b>...</b>	...	...	...	...
<b>164</b>	False	False	False	False
<b>165</b>	False	False	False	False

	Duration	Pulse	Maxpulse	Calories
166	False	False	False	False
167	False	False	False	False
168	False	False	False	False

169 rows × 4 columns

In [5]:

```
#checking is there any null value is there or not.
df.isnull().values.any()
```

Out[5]:

True

In [7]:

```
# a. Replace the null values with the mean
new_df=df.fillna(df.mean())
```

In [10]:

```
new_df.isnull().values.any()
```

Out[10]:

False

In [11]:

```
#4. Select at least two columns and aggregate the data using: min, max,
count, mean

# by using groupby function with aggregation to get mean, min and max
values
```

```

result = df.groupby('Duration').agg({'Calories': ['mean', 'min', 'max']})

print("Mean, min, and max values are")
print(result)

```

```

Mean, min, and max values are
      Calories
      mean    min    max
Duration
15      87.350000   50.5  124.2
20     151.600000   50.3  229.4
25     244.200000  244.2  244.2
30     192.125000   86.2  319.2
45     273.236364  100.7  406.0
60     339.675000  215.2  486.0
75     325.400000  320.4  330.4
80     643.100000  643.1  643.1
90     541.800000  466.4  700.0
120     666.833333  500.0 1000.1
150     939.400000  816.0 1115.0
160     943.700000  853.0 1034.4
180     733.600000  600.1  800.4
210    1618.200000 1376.0 1860.4
270    1729.000000 1729.0 1729.0
300    1500.200000 1500.2 1500.2

```

In [12]:

```

#5. Filter the dataframe to select the rows with calories values between
500 and 1000.

```

```

df.query('Calories < 1000 and Calories > 500')

```

Out[12]:

	Duration	Pulse	Maxpulse	Calories
<b>51</b>	80	123	146	643.1
<b>62</b>	160	109	135	853.0
<b>65</b>	180	90	130	800.4

	Duration	Pulse	Maxpulse	Calories
66	150	105	135	873.4
67	150	107	130	816.0
72	90	100	127	700.0
73	150	97	127	953.2
75	90	98	125	563.2
78	120	100	130	500.4
90	180	101	127	600.1
99	90	93	124	604.1
103	90	90	100	500.4
106	180	90	120	800.3
108	90	90	120	500.3

In [13]:

```
# 6. Filter the dataframe to select the rows with calories values > 500
and pulse < 100
df.query('Calories > 500 and Pulse < 100')
```

Out[13]:

	Duration	Pulse	Maxpulse	Calories
65	180	90	130	800.4

	Duration	Pulse	Maxpulse	Calories
70	150	97	129	1115.0
73	150	97	127	953.2
75	90	98	125	563.2
99	90	93	124	604.1
103	90	90	100	500.4
106	180	90	120	800.3
108	90	90	120	500.3

In [14]:

```
#7. Create a new "df_modified" dataframe that contains all the columns  
from df except for "Maxpulse"
df_modified=df.drop(columns=["Maxpulse"])
df_modified
```

Out[14]:

	Duration	Pulse	Calories
0	60	110	409.1
1	60	117	479.0
2	60	103	340.0
3	45	109	282.4

	Duration	Pulse	Calories
<b>4</b>	45	117	406.0
...	...	...	...
<b>164</b>	60	105	290.8
<b>165</b>	60	110	300.0
<b>166</b>	60	115	310.2
<b>167</b>	75	120	320.4
<b>168</b>	75	125	330.4

169 rows × 3 columns

In [17]:

```
# 8. Delete the "Maxpulse" column from the main df dataframe
df.drop(columns=["Maxpulse"], axis=1, inplace=True)
df
```

Out[17]:

	Duration	Pulse	Calories
<b>0</b>	60	110	409.1
<b>1</b>	60	117	479.0
<b>2</b>	60	103	340.0
<b>3</b>	45	109	282.4



	Duration	Pulse	Calories
<b>4</b>	45	117	406.0
...	...	...	...
<b>164</b>	60	105	290.8
<b>165</b>	60	110	300.0
<b>166</b>	60	115	310.2
<b>167</b>	75	120	320.4
<b>168</b>	75	125	330.4

169 rows × 3 columns

In [22]:

```
#9. Convert the datatype of Calories column to int datatype.
df=df.fillna(df.mean())
df = df.astype({'Calories':'int'})

print(df.dtypes)
```

```
Duration      int64
Pulse         int64
Calories      int64
dtype: object
```

In [23]:

```
#Using pandas create a scatter plot for the two columns (Duration and Calories).


df.plot(kind = 'scatter', x = 'Duration', y = 'Calories')
```

```
/usr/local/lib/python3.9/dist-packages/pandas/plotting/_matplotlib/core.py:11
14: UserWarning: No data for colormapping provided via 'c'. Parameters 'cmap'
will be ignored
```

```
scatter = ax.scatter(
```

Out[23]:

```
<Axes: xlabel='Duration', ylabel='Calories'>
```



## 1. (Titanic Dataset)

1. Find the correlation between 'survived' (target column) and 'sex' column for the Titanic use case inclass.

a. Do you think we should keep this feature?

2. Do at least two visualizations to describe or show correlations.

3. Implement Naïve Bayes method using scikit-learn library and report the accuracy.

In [24]:

```
#import data
```

```
test_df = pd.read_csv("/content/drive/MyDrive/Colab
Notebooks/Dataset/test.csv")

train_df = pd.read_csv("/content/drive/MyDrive/Colab
Notebooks/Dataset/train.csv")
```

In [25]:

```
#Data ANalysis
```

```
train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   PassengerId     891 non-null   int64
 1   Survived        891 non-null   int64
 2   Pclass          891 non-null   int64
 3   Name            891 non-null   object
 4   Sex             891 non-null   object
 5   Age             714 non-null   float64
 6   SibSp           891 non-null   int64
 7   Parch           891 non-null   int64
 8   Ticket          891 non-null   object
 9   Fare            891 non-null   float64
10   Cabin           204 non-null   object
11   Embarked        889 non-null   object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

In [26]:

```
train_df.describe()
```

Out[26]:

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

In [27]:

```
import seaborn as sns
%matplotlib inline
from matplotlib import pyplot as plt

survived = 'Survived'
not_survived = 'not survived'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(10, 4))
women = train_df[train_df['Sex']=='female']
men = train_df[train_df['Sex']=='male']
ax = sns.distplot(women[women[survived]==1].Age.dropna(), bins=18, label = survived, ax = axes[0], kde =False)
ax = sns.distplot(women[women[survived]==0].Age.dropna(), bins=40, label = not_survived, ax = axes[0], kde =False)
ax.legend()
ax.set_title('Female')
ax = sns.distplot(men[men[survived]==1].Age.dropna(), bins=18, label = survived, ax = axes[1], kde = False)
```

```
ax = sns.distplot(men[men['Survived']==0].Age.dropna(), bins=40, label =
not_survived, ax = axes[1], kde = False)

ax.legend()

_ = ax.set_title('Male')
```

<ipython-input-27-0cf8acbf0d6>:10: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
ax = sns.distplot(women[women['Survived']==1].Age.dropna(), bins=18, label
= survived, ax = axes[0], kde =False)
```

<ipython-input-27-0cf8acbf0d6>:11: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
ax = sns.distplot(women[women['Survived']==0].Age.dropna(), bins=40, label
= not_survived, ax = axes[0], kde =False)
```

<ipython-input-27-0cf8acbf0d6>:14: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
ax = sns.distplot(men[men['Survived']==1].Age.dropna(), bins=18, label = su
rvived, ax = axes[1], kde = False)
```


<ipython-input-27-0cf8acbf0d6>:15: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
ax = sns.distplot(men[men['Survived']==0].Age.dropna(), bins=40, label = not_survived, ax = axes[1], kde = False)
```



In [30]:

```
from matplotlib import pyplot as plt
import pandas as pd
import numpy as np
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score

train_data = pd.read_csv('/content/drive/MyDrive/Colab
Notebooks/Dataset/train.csv')

test_data = pd.read_csv('/content/drive/MyDrive/Colab
Notebooks/Dataset/test.csv')

# Find the correlation between Survived and Sex

corr =
train_data['Survived'].corr(train_data['Sex'].astype('category').cat.codes
)

print("Correlation between Survived and Sex: ",corr)
```

```
print('a. Do you think we should keep this feature?')


print('Yes, we should keep this feature as it has a correlation of', corr,
      'with the target variable but some other features can be dropped as they
      have very less correlation with the target variable.')

# Do at least two visualizations to describe the data
# Histogram of age
train_data['Age'].plot.hist(title='Histogram of Age')
plt.show()

# Scatter plot of age and fare
train_data.plot.scatter(x='Age', y='Fare', title='Scatter plot of Age and
Fare')
plt.show()

# Plot between age and survived
train_data.plot.scatter(x='Age', y='Survived', title='Scatter plot of Age
and Survived')
plt.show()
```


```
Correlation between Survived and Sex: -0.5433513806577555
a. Do you think we should keep this feature?
Yes, we should keep this feature as it has a correlation of -0.54335138065775
55 with the target variable but some other features can be dropped as they ha
ve very less correlation with the target variable.
```

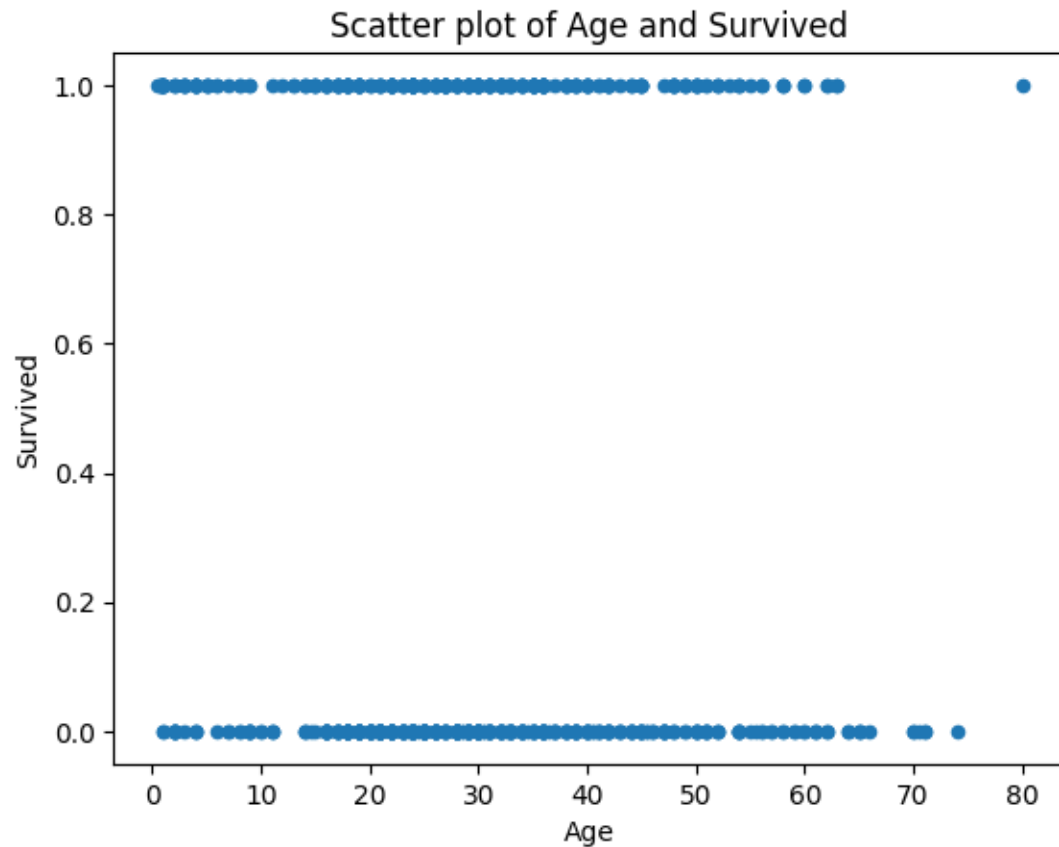


```
/usr/local/lib/python3.9/dist-packages/pandas/plotting/_matplotlib/core.py:11
14: UserWarning: No data for colormapping provided via 'c'. Parameters 'cmap'
will be ignored
    scatter = ax.scatter(
```



Scatter plot of Age and Fare





## 2. (Glass Dataset)

1. Implement Naïve Bayes method using scikit-learn library. a. Use the glass dataset available in [Link](#) also provided in your assignment. b. Use `train_test_split` to create training and testing part.
2. Evaluate the model on testing part using score and Do at least two visualizations to describe or show correlations in the Glass Dataset.

In [32]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, accuracy_score
import warnings
```

```
warnings.filterwarnings("ignore")

# a. read glass.csv file as a dataframe
glass_data = pd.read_csv('/content/drive/MyDrive/Colab
Notebooks/Dataset/glass.csv')

# b. Use train_test_split to create training and testing part
# Split the data into training and testing data
X_train, X_test, y_train, y_test =
train_test_split(glass_data.drop('Type', axis=1), glass_data['Type'],
test_size=0.3, random_state=42)

# implement Naïve Bayes method using scikit-learn library and Evaluate the
model on testing part using score and classification_report

# Create a Gaussian Classifier
model = GaussianNB()

# Train the model using the training sets
model.fit(X_train, y_train)

# Predict the response for test dataset
y_pred = model.predict(X_test)

# Calculate the accuracy of the model
print("Accuracy of the Naive Bayes model: ", accuracy_score(y_test,
y_pred))

# Print classification report
print('Classification Report for Naive Bayes model: ')
print(classification_report(y_test, y_pred))
```

```
# Use SVM method using scikit-learn library and Evaluate the model on
testing part using score and classification_report

# Create a SVM Classifier
model = SVC(kernel='linear')

# Train the model using the training sets
model.fit(X_train, y_train)

# Predict the response for test dataset
y_pred = model.predict(X_test)

# Calculate the accuracy of the model
print("Accuracy of the SVM model: ", accuracy_score(y_test, y_pred))

# Print classification report
print('Classification Report for SVM model: ')
print(classification_report(y_test, y_pred))

# Do at least two visualizations to describe or show correlations in the
Glass Dataset

# Histogram of refractive index
glass_data['RI'].plot.hist(title='Histogram of Refractive Index')
plt.show()

# Scatter plot of refractive index and Ca
glass_data.plot.scatter(x='RI', y='Ca', title='Scatter plot of Refractive
Index and Ca')
plt.show()

print('Which algorithm you got better accuracy? Can you justify why?')
# accuracy of Naive Bayes model:  0.3076923076923077
```

```
# accuracy of SVM model: 0.676923076923077
```

```
print('SVM model has better accuracy than Naive Bayes model. This is  
because SVM model tries to find the best possible decision boundary  
between the data points of different classes. It tries to maximize the  
margin between the decision boundary and the data points. On the other  
hand, Naive Bayes model assumes that the features are independent of each  
other and tries to find the probability of the data point belonging to a  
particular class. Hence, SVM model has better accuracy than Naive Bayes  
model.')
```

Accuracy of the Naive Bayes model: 0.3076923076923077

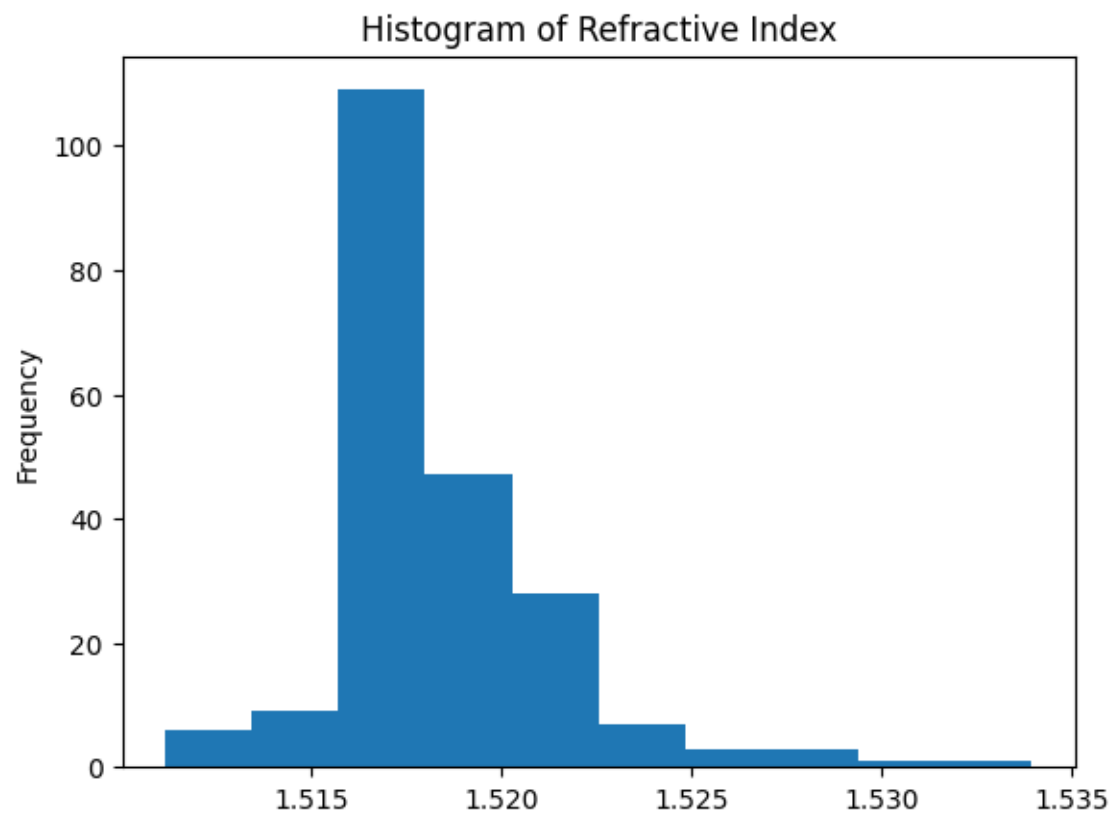
Classification Report for Naive Bayes model:


	precision	recall	f1-score	support
1	0.00	0.00	0.00	19
2	0.40	0.17	0.24	23
3	0.08	0.75	0.15	4
5	0.33	0.17	0.22	6
6	0.75	1.00	0.86	3
7	0.90	0.90	0.90	10
accuracy			0.31	65
macro avg	0.41	0.50	0.40	65
weighted avg	0.35	0.31	0.29	65

Accuracy of the SVM model: 0.676923076923077

Classification Report for SVM model:

	precision	recall	f1-score	support
1	0.65	0.79	0.71	19
2	0.59	0.70	0.64	23
3	0.00	0.00	0.00	4
5	0.75	0.50	0.60	6
6	0.50	0.33	0.40	3
7	1.00	0.90	0.95	10
accuracy			0.68	65
macro avg	0.58	0.54	0.55	65
weighted avg	0.65	0.68	0.65	65





Which algorithm you got better accuracy? Can you justify why?

SVM model has better accuracy than Naive Bayes model. This is because SVM model tries to find the best possible decision boundary between the data points of different classes. It tries to maximize the margin between the decision boundary and the data points. On the other hand, Naive Bayes model assumes that the features are independent of each other and tries to find the probability of the data point belonging to a particular class. Hence, SVM model has better accuracy than Naive Bayes model.