

Spring 2023: CS5710 – Machine Learning

In-Class Programming Assignment-5

GitHub Link - https://github.com/raimukul/MachineLearning_Assignments/tree/main/Assignment%2005

Video link-

https://drive.google.com/file/d/13cmfNAbHYzUXgphNMItYNpi80udNLryJ/view?usp=share_link

Code:

1. Principal Component Analysis

- a. Apply PCA on CC dataset.
- b. Apply k-means algorithm on the PCA result and report your observation if the silhouette score has improved or not?
- c. Perform Scaling+PCA+K-Means and report performance.

In [1]:

```
import pandas as pd
import numpy as np
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.model_selection import train_test_split
from sklearn.metrics import silhouette_score
import warnings
warnings.filterwarnings('ignore')

# read dataset
df = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/Assignment 05/datasets/CC GENERAL.csv')
# drop CUST_ID column
```

```
df.drop('CUST_ID', axis=1, inplace=True)

# drop rows with missing values

df.dropna(inplace=True)


# split dataset into train and test

X_train, X_test = train_test_split(df, test_size=0.2, random_state=42)


# scale fit training data

scaler = StandardScaler()

scaler.fit(X_train)


# apply transform to training and test data

X_train = scaler.transform(X_train)

X_test = scaler.transform(X_test)


# Apply k-means algorithm on the original data

kmeans = KMeans(n_clusters=2, random_state=42)

kmeans.fit(X_train)

y_pred = kmeans.predict(X_train)

sil_original = silhouette_score(X_train, y_pred)

print('Silhouette score for k-means on original data: ', sil_original)


# apply PCA to training and test data

pca = PCA(n_components=2)

pca.fit(X_train)

X_train = pca.transform(X_train)

X_test = pca.transform(X_test)


kmeans = KMeans(n_clusters=2, random_state=42)

kmeans.fit(X_train)
```

```

y_pred = kmeans.predict(X_train)
sil_pca = silhouette_score(X_train, y_pred)
print('Silhouette score for k-means on PCA result: ', sil_pca)

print('Silhouette score for k-means on original data is ', sil_original, ' and silhouette score for k-means on
PCA result is ', sil_pca)
if(sil_pca > sil_original):
    print('Silhouette score has improved')
else:
    print('Silhouette score has not improved')

# report performance on test data
y_pred = kmeans.predict(X_test)
sil_test = silhouette_score(X_test, y_pred)
print('Silhouette score for k-means on test data: ', sil_test)

```

Silhouette score for k-means on original data: 0.21163643243769295
 Silhouette score for k-means on PCA result: 0.46015232772144554
 Silhouette score for k-means on original data is 0.21163643243769295 and silhouette score for k-means
 on PCA result is 0.46015232772144554
 Silhouette score has improved
 Silhouette score for k-means on test data: 0.460742753941452

```
sil_pca = silhouette_score(X_train, y_pred)
print('silhouette score for k-means on PCA result: ', sil_pca)

print('silhouette score for k-means on original data is ', sil_original, ' and silhouette score for k-means on PCA result is ', sil_pca)
if(sil_pca > sil_original):
    print('silhouette score has improved')
else:
    print('silhouette score has not improved')

# report performance on test data
y_pred = kmeans.predict(X_test)
sil_test = silhouette_score(X_test, y_pred)
print('silhouette score for k-means on test data: ', sil_test)
```

silhouette score for k-means on original data: 0.21163643243769295
silhouette score for k-means on PCA result: 0.46015232772144554
silhouette score for k-means on original data is 0.21163643243769295 and silhouette score for k-means on PCA result is 0.46015232772144554
silhouette score has improved
silhouette score for k-means on test data: 0.460742753941452

2. Use pd_speech_features.csv

- a. Perform Scaling
- b. Apply PCA (k=3)
- c. Use SVM to report performance

2. Use pd_speech_features.csv

a. Perform Scaling

b. Apply PCA (k=3)

c. Use SVM to report performance

In [6]:

```
import pandas as pd
```

```
import numpy as np
```

```
from sklearn.cluster import KMeans
```

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.decomposition import PCA
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.metrics import silhouette_score
```

```
import warnings
```

```
warnings.filterwarnings('ignore')
```

```
from sklearn.svm import SVC
```

```
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report

# Use pd_speech_features.csv
df = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/Assignment
05/datasets/pd_speech_features.csv')
# drop id column
df.drop('id', axis=1, inplace=True)
# drop rows with missing values
df.dropna(inplace=True)

X = df.drop('class', axis=1)
y = df['class']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# a. Perform Scaling
scaler = StandardScaler()
scaler.fit(X_train)

# apply transform to training and test data
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)

# b. Apply PCA (k=3)
pca = PCA(n_components=3)
pca.fit(X_train)
X_train = pca.transform(X_train)
X_test = pca.transform(X_test)
```

c. Use SVM to report performance

```
svm = SVC()
```

```
svm.fit(X_train, y_train)
```

```
y_pred = svm.predict(X_test)
```

```
print('Accuracy score: ', accuracy_score(y_test, y_pred))
```

```
print('Confusion matrix: ', confusion_matrix(y_test, y_pred))
```

```
print('Classification report: ', classification_report(y_test, y_pred))
```

Accuracy score: 0.8026315789473685

Confusion matrix: [[16 22]

[8 106]]

Classification report: precision recall f1-score support

0	0.67	0.42	0.52	38
---	------	------	------	----

1	0.83	0.93	0.88	114
---	------	------	------	-----

accuracy		0.80	152
----------	--	------	-----

macro avg	0.75	0.68	0.70	152
-----------	------	------	------	-----

weighted avg	0.79	0.80	0.79	152
--------------	------	------	------	-----

```
# apply transform to training and test data
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)

# b. Apply PCA (k=3)
pca = PCA(n_components=3)
pca.fit(X_train)
X_train = pca.transform(X_train)
X_test = pca.transform(X_test)

# c. Use SVM to report performance
svm = SVC()
svm.fit(X_train, y_train)
y_pred = svm.predict(X_test)
print('Accuracy score: ', accuracy_score(y_test, y_pred))
print('Confusion matrix: ', confusion_matrix(y_test, y_pred))
print('Classification report: ', classification_report(y_test, y_pred))
```

Accuracy score: 0.8026315789473685
Confusion matrix: [[16 22]
[8 106]]
Classification report:

		precision	recall	f1-score	support
0	0.67	0.42	0.52	38	
1	0.83	0.93	0.88	114	
accuracy			0.80	152	
macro avg	0.75	0.68	0.70	152	
weighted avg	0.79	0.80	0.79	152	

3. Apply Linear Discriminant Analysis (LDA) on Iris.csv dataset to reduce dimensionality of data tok=2.

In [5]:

```
import pandas as pd
```

```
import numpy as np
```

```
# import lda
```

```
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
```

```
# read dataset
```

```
df = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/Assignment 05/datasets/Iris.csv')
```

```
# drop id column
```

```
df.drop('Id', axis=1, inplace=True)
```

```
# drop rows with missing values
```

```
df.dropna(inplace=True)
```

```
# split dataset into train and test
```

```
X = df.drop('Species', axis=1)
```

```
y = df['Species']
```

```
# apply LDA to training and test data
```

```
lda = LinearDiscriminantAnalysis(n_components=2)
```

```
lda.fit(X, y)
```

```
X = lda.transform(X)
```

```
print(X)
```

```
[[ 8.0849532  0.32845422]
 [ 7.1471629 -0.75547326]
 [ 7.51137789 -0.23807832]
 [ 6.83767561 -0.64288476]
 [ 8.15781367  0.54063935]
 [ 7.72363087  1.48232345]
 [ 7.23514662  0.3771537 ]
 [ 7.62974497  0.01667246]
 [ 6.58274132 -0.98737424]
 [ 7.36884116 -0.91362729]
 [ 8.42181434  0.67622968]
 [ 7.24739721 -0.08292417]
 [ 7.35062105 -1.0393597 ]
 [ 7.59646896 -0.77671553]
 [ 9.86936588  1.61486093]
 [ 9.18033614  2.75558626]
 [ 8.59760709  1.85442217]
 [ 7.7995682  0.60905468]
 [ 8.1000091  0.99610981]
 [ 8.04543611  1.16244332]
 [ 7.52046427 -0.156233 ]
 [ 7.60526378  1.22757267]
 [ 8.70408249  0.89959416]
 [ 6.26374139  0.46023935]
 [ 6.59191505 -0.36199821]
 [ 6.79210164 -0.93823664]
 [ 6.84048091  0.4848487 ]
 [ 7.948386   0.23871551]
 [ 8.01209273  0.11626909]
 [ 6.85589572 -0.51715236]
 [ 6.78303525 -0.72933749]
 [ 7.38668238  0.59101728]
 [ 9.16249492  1.25094169]
 [ 9.49617185  1.84989586]
```


[7.36884116 -0.91362729]
[7.9756525 -0.13519572]
[8.63115466 0.4346228]
[7.36884116 -0.91362729]
[6.95602269 -0.67887846]
[7.71167183 0.01995843]
[7.9361354 0.69879338]
[5.6690533 -1.90328976]
[7.26559733 -0.24793625]
[6.42449823 1.26152073]
[6.88607488 1.07094506]
[6.77985104 -0.47815878]
[8.11232705 0.78881818]
[7.21095698 -0.33438897]
[8.33988749 0.6729437]
[7.69345171 -0.10577397]
[-1.45772244 0.04186554]
[-1.79768044 0.48879951]
[-2.41680973 -0.08234044]
[-2.26486771 -1.57609174]
[-2.55339693 -0.46282362]
[-2.41954768 -0.95728766]
[-2.44719309 0.79553574]
[-0.2160281 -1.57096512]
[-1.74591275 -0.80526746]
[-1.95838993 -0.35044011]
[-1.19023864 -2.61561292]
[-1.86140718 0.32050146]
[-1.15386577 -2.61693435]
[-2.65942607 -0.63412155]
[-0.38024071 0.09211958]
[-1.20280815 0.09561055]
[-2.7626699 0.03156949]
[-0.76227692 -1.63917546]
[-3.50940735 -1.6724835]
[-1.08410216 -1.6100398]
[-3.71895188 1.03509697]
[-0.99937 -0.47902036]
[-3.83709476 -1.39488292]
[-2.24344339 -1.41079358]
[-1.25428429 -0.53276537]
[-1.43952232 -0.12314653]
[-2.45921948 -0.91961551]
[-3.52471481 0.16379275]
[-2.58974981 -0.17075771]
[0.31197324 -1.29978446]
[-1.10232227 -1.7357722]
[-0.59844322 -1.92334798]
[-0.89605882 -0.89192518]
[-4.49567379 -0.87924754]
[-2.9265236 0.02499754]

[-2.10119821 1.18719828]
[-2.14367532 0.09713697]
[-2.48342912 -1.92190266]
[-1.31792367 -0.15753271]
[-1.95529307 -1.14514953]
[-2.38909697 -1.5823776]
[-2.28614469 -0.32562577]
[-1.26934019 -1.20042096]
[-0.28888857 -1.78315025]
[-2.00077969 -0.8969707]
[-1.16910587 -0.52787187]
[-1.6092782 -0.46274252]
[-1.41813799 -0.53933732]
[0.47271009 -0.78924756]
[-1.54557146 -0.58518894]
[-7.85608083 2.11161905]
[-5.5156825 -0.04401811]
[-6.30499392 0.46211638]
[-5.60355888 -0.34236987]
[-6.86344597 0.81602566]
[-7.42481805 -0.1726265]
[-4.68086447 -0.50758694]
[-6.31374875 -0.96068288]
[-6.33198886 -1.37715975]
[-6.87287126 2.69458147]
[-4.45364294 1.33693971]
[-5.4611095 -0.21035161]
[-5.67679825 0.82435717]
[-5.97407494 -0.10462115]
[-6.78782019 1.5744553]
[-5.82871291 1.98940576]
[-5.0664238 -0.02730214]
[-6.60847169 1.7420041]
[-9.18829265 -0.74909806]
[-4.76573133 -2.14417884]
[-6.29305487 1.63373692]
[-5.37314577 0.63153087]
[-7.58557489 -0.97390788]
[-4.38367513 -0.12213933]
[-5.73135125 1.28143515]
[-5.27583147 -0.0384815]
[-4.0923206 0.18307048]
[-4.08316687 0.51770204]
[-6.53257435 0.28724638]
[-4.577648 -0.84457527]
[-6.23500611 -0.70621819]
[-5.21836582 1.46644917]
[-6.81795935 0.56784684]
[-3.80972091 -0.93451896]
[-5.09023453 -2.11775698]
[-6.82119092 0.85698379]

```

[-6.54193229 2.41858841]
[-4.99356333 0.18488299]
[-3.94659967 0.60744074]
[-5.22159002 1.13613893]
[-6.67858684 1.785319 ]
[-5.13687786 1.97641389]
[-5.5156825 -0.04401811]
[-6.81196984 1.44440158]
[-6.87289126 2.40383699]
[-5.67401294 1.66134615]
[-5.19712883 -0.36550576]
[-4.98171163 0.81297282]
[-5.90148603 2.32075134]
[-4.68400868 0.32508073]]

```

Outputs:

The screenshot shows a Jupyter Notebook titled "Assignment 5.ipynb" in a web browser. The code cell contains the following Python code:

```

X = lda.transform(X)
print(X)

```

The output of the code is a 2D array of numerical values, displayed as a list of lists. The array has 20 rows and 2 columns. The values are as follows:

```

[[ 8.0849532  0.32845422]
 [ 7.1471629 -0.75547326]
 [ 7.51137789 -0.23807832]
 [ 6.83767561 -0.64288476]
 [ 8.15781367  0.54063935]
 [ 7.72363087  1.48232345]
 [ 7.23514662  0.3771537 ]
 [ 7.62974497  0.01667246]
 [ 6.58274132 -0.98737424]
 [ 7.36884116 -0.91362729]
 [ 8.42181434  0.67522968]
 [ 7.24739721 -0.08292417]
 [ 7.35062105 -1.0393597 ]
 [ 7.59646896 -0.77671553]
 [ 9.86936588  1.61486093]
 [ 9.18033614  2.75558626]
 [ 8.59760709  1.85442217]
 [ 7.7995682  0.60905468]
 [ 8.1000091  0.99610981]
 [ 8.04543611  1.16244332]
 [ 7.52046427 -0.156233 ]
 [ 7.60526378  1.22757267]
 [ 8.70408249  0.89959416]
 [ 6.26374139  0.46023935]
 [ 6.59191505 -0.36199821]
 [ 6.79210164 -0.93823664]
 [ 6.84048091  0.4848487 ]
 [ 7.948386  0.23871551]
 [ 8.01209273  0.11626909]
 [ 6.85589572 -0.51715236]]

```

The notebook interface includes a menu bar (File, Edit, View, Insert, Runtime, Tools, Help), a toolbar with icons for code, text, and search, and a status bar at the bottom indicating the completion time as 9:52 PM.

```
X = lda.transform(X)

print(X)
```

```
[[-1.45772244  0.04186554]
 [-1.79768044  0.48879951]
 [-2.41680973 -0.08234044]
 [-2.26486771 -1.57609174]
 [-2.55339693 -0.46282362]
 [-2.41954768 -0.95728766]
 [-2.44719309  0.79535374]
 [-0.2160281  -1.57096512]
 [-1.74591275 -0.80526746]
 [-1.95838993 -0.35044011]
 [-1.19023864 -2.61561292]
 [-1.86140718  0.32050146]
 [-1.15386577 -2.61693435]
 [-2.65942607 -0.63412155]
 [-0.38024071  0.09211958]
 [-1.20280815  0.09561055]
 [-2.7626699  0.03156049]
 [-0.76227692 -1.63917546]
 [-3.50940735 -1.6724835 ]
 [-1.08410216 -1.6100398 ]
 [-3.71895188  1.03509697]
 [-0.99937    -0.47902036]
 [-3.83709476 -1.39488292]
 [-2.24344339 -1.41079358]
 [-1.25428429 -0.53276537]
 [-1.43952232  0.12314653]
 [-2.45921948 -0.91961551]
 [-3.52471481  0.16379275]
 [-2.58974981 -0.17075771]
 [ 0.31197324 -1.29978446]]
```

```
X = lda.transform(X)

print(X)
```

```
[[-2.48342912 -1.92190266]
 [-1.31792367 -0.15753271]
 [-1.95529307 -1.14514953]
 [-2.38909697 -1.5823776 ]
 [-2.28614469 -0.32562577]
 [-1.26934019 -1.20042096]
 [-0.2888857 -1.78315025]
 [-2.00077969 -0.8969707 ]
 [-1.16910587 -0.52787187]
 [-1.6092782  -0.46274252]
 [-1.41813799 -0.53933732]
 [ 0.47271009 -0.78924756]
 [-1.54557146 -0.58518894]
 [-7.85608083  2.11161905]
 [-5.5156825  -0.04401811]
 [-6.30499392  0.46211638]
 [-5.60355888 -0.34236987]
 [-6.86344597  0.81602566]
 [-7.42481805 -0.1726265 ]
 [-4.68086447 -0.50750694]
 [-6.31374875 -0.96068288]
 [-6.33198886 -1.37715975]
 [-6.87287126  2.69458147]
 [-4.45364294  1.33693971]
 [-5.4611095  -0.21035161]
 [-5.67679825  0.82435717]
 [-5.97407494 -0.10462115]
 [-6.78782019  1.5744553 ]
 [-5.82871291  1.98940576]
 [-5.0664238  -0.02730214]]
```

4. Briefly identify the difference between PCA and LDA

Answer - PCA is an unsupervised algorithm that is used to reduce the dimensionality of the data. It is used to find the principal components of the data. LDA is also a supervised algorithm that is used to reduce the dimensionality of the data. It is used to find the linear combination of features that characterizes or separates two or more classes.

